

## Popular Models

| Company   | Latest Model (Private) | Release Date | Knowledge Cutoff | Cost per 1M Tokens            | Context Window     | Public / Private      | Best Use Case                                     | Latest Public Model (if applicable)                    |
|-----------|------------------------|--------------|------------------|-------------------------------|--------------------|-----------------------|---------------------------------------------------|--------------------------------------------------------|
| OpenAI    | GPT-5.5                | Apr-26       | Nov-25           | ~\$2.50 input / \$10 output   | 400K (128K output) | Private               | Complex reasoning, agentic workflows              | GPT-4 Turbo (Nov 2023) — still the latest open weights |
| Anthropic | Claude Opus 4.7        | Apr-26       | Jan-26           | \$3 input / \$15 output       | 200K (64K output)  | Private               | Precision coding, structured workflows            | No public release (Anthropic is API-only)              |
| Google    | Gemini 3 Pro           | Apr-26       | Jan-26           | \$0.15 input / \$0.60 output  | 1M (8K output)     | Private               | Multimodal (text+image+video+audio)               | No public release (Google is API-only)                 |
| Mistral   | Large 2                | Mar-25       | Late 2024        | ~\$0.07 input / \$0.28 output | 128K               | Public (open weights) | Multilingual reasoning, cost-efficient deployment | Same model (open weights)                              |
| Meta      | LLaMA 3.3              | Apr-25       | Early 2025       | Free (open weights)           | 10M (MoE)          | Public (open weights) | Custom fine-tuning, research                      | Same model (open weights)                              |
| Alibaba   | Qwen 2.5               | Jan-25       | Jun-24           | ~\$0.36 input / \$0.40 output | 128K               | Public (open weights) | Chinese NLP, balanced English                     | Same model (open weights)                              |
| DeepSeek  | V3                     | Feb-26       | Sep-25           | ~\$0.27 input / \$1.10 output | 1M                 | Public (open weights) | Math, coding, cost efficiency                     | Same model (open weights)                              |

|            |                     |        |        |                               |      |         |                                         |                                     |
|------------|---------------------|--------|--------|-------------------------------|------|---------|-----------------------------------------|-------------------------------------|
| xAI        | Grok 4.3            | Apr-26 | Dec-25 | ~\$1.25 input / \$2.50 output | 1M   | Private | Conversational reasoning, X integration | No public release (xAI is API-only) |
| Baidu      | ERNIE Bot 4.5 Turbo | Apr-25 | Mar-25 | ~\$0.07 input / \$0.28 output | 120K | Private | Chinese multimodal, Baidu ecosystem     | No public release                   |
| Inflection | Pi (Inflection 3)   | Oct-25 | Oct-25 | ~\$2.50 input / \$10 output   | 8K   | Private | Personal assistant, empathetic dialogue | No public release                   |

### Less Popular Models

| Company            | Latest Model                 | Release Date     | Knowledge Cutoff | Cost per 1M Tokens         | Context Window    | Public / Private      | Best Use Case                                  |
|--------------------|------------------------------|------------------|------------------|----------------------------|-------------------|-----------------------|------------------------------------------------|
| <b>AI21 Labs</b>   | Jurassic-2 Ultra (update)    | Early 2026       | Mid-2022         | ~\$18.80 input/output      | 8K                | Private               | Long-form writing, enterprise text generation  |
| <b>Cohere</b>      | Command R+                   | Aug-25           | Feb-23           | ~\$2.5 input / \$10 output | 128K              | Private               | Retrieval-augmented generation (RAG)           |
| <b>Aleph Alpha</b> | Luminous v5.5                | 2025             | N/A              | ~\$30 input / \$33 output  | ~32K              | Private               | Multilingual + explainability, transparency    |
| <b>Huawei</b>      | PanGu 5.5                    | Jun-25           | N/A              | Enterprise pricing         | Large multimodal  | Private               | Scientific modeling, physics simulations       |
| <b>TII (UAE)</b>   | Falcon 180B (still relevant) | Sep 2025 update  | Mid-2022         | Free (Apache 2.0)          | ~32K              | Public (open weights) | Large-scale open-source research               |
| <b>01.AI</b>       | Yi-34B (latest variant)      | Late 2025 update | N/A              | ~\$0.90 input/output       | 32K–200K variants | Public (open weights) | Balanced reasoning + coding, regional adoption |